

Agent Eval Mutation Lab

Deep Research, GPT-5.6 Pro decision review, prior-art correction, and verified offline Python milestone

Decision: build a narrow scorer-mutation benchmark, not a copied framework, authorization-gate clone, or another consumer app.

Executive answer

The current resume proves unusually broad product engineering, but Python is visible only in the skills list. Current target roles repeatedly combine production Python with evaluation, agent tooling, measurement, open source, or empirical work. The highest-leverage response is one original public Python artifact with a reproducible empirical question - not a fork-and-rename portfolio clone.

Two GPT-5.6 Pro reviews selected an offline, framework-independent benchmark architecture and then approved the concrete scorer-mutation subject after an earlier argument-aware authorization idea failed the prior-art gate. The selected project tests whether scorers preserve distinctions among proposed calls, actual execution, transient harm, final state, and missing receipts.

What is verified today

- 13 synthetic cases across five scenario families and seven execution-semantic mutation types.
- Three substantive scorer contracts plus always-safe, always-attack, and always-abstain controls.
- Sixteen focused tests, Ruff, and strict mypy all pass.
- The run is offline, needs no API key, and emits deterministic JSON and Markdown reports.
- A fail-closed Inspect AI adapter was verified against fresh full JSON logs, and receipt ablations expose a concrete v1 scorer weakness.
- No real-model, framework-safety, production-safety, or independent-fluency claim is supported yet.

1. Evidence and role alignment

Verified evidence	Implication
Current named projects are TypeScript, JavaScript, or Deno.	Python is asserted in Skills but lacks a named flagship artifact.
Anthropic, Scale, Quadrillion, and agent-platform roles name Python, empirical work, evals, or agent tooling.	A Python evaluation artifact closes a cross-role gap rather than a single-posting keyword gap.
Inspect AI, ControlArena, and AgentDojo are active MIT-licensed Python ecosystems.	Use them later as dependencies or adapters; do not fork and lightly restyle them.
Agent Security Gate already covers argument-aware enforcement and related benchmarks.	The first project idea was rejected as overlapping prior art.
AgentDojo issue #168 documents attempted-but-blocked calls scored as success.	Motivates scorer validation, but the issue and reproduction are not an original discovery.

Role-source highlights

[Anthropic Fellows](#) names strong Python, empirical AI-safety research, independent execution, open source, and clear communication. [Scale AI Builder Intern](#) names production-quality Python and/or JavaScript, agentic workflows, LLM evals, measurement, and portfolio evidence. [Quadrillion](#) asks for significant engineering ability in Python or React while building an agentic research platform.

Framework reuse decision

[Inspect AI](#) supplies task, dataset, solver, scorer, and log abstractions. [ControlArena](#) supplies AI-control settings, protocols, monitors, scorers, and safety/usefulness analysis. [AgentDojo](#) supplies realistic prompt-injection environments and programmatic scoring. All are strong later adapter targets. None should be copied wholesale merely to produce a resume project.

2. Prior-art correction and research question

The first proposed subject - argument-aware pre-execution authorization - was rejected after finding [Agent Security Gate](#), an Apache-2.0 Python project with policy enforcement, approvals, audit receipts, an 18-scenario benchmark, and an AgentDojo study.

The revised subject is broader than reproducing [AgentDojo issue #168](#): **Can execution-semantic mutation testing expose tool-agent scorers that confuse proposed actions, actual execution, and realized harm?**

Honest originality boundary

Mutation testing is established in [SWE-Mutation](#), [LLMorpheus](#), and recent semantic-evaluator work such as [Breaking Models to Test the Judge](#). The defensible claim is an adaptation to execution-semantic tool-agent scorers with an explicit outcome ontology and finite benchmark. It is not a first-ever mutation-testing claim or discovery of the AgentDojo bug.

3. Benchmark design

Outcome ontology

- Unsafe attempt: a prohibited tool call was proposed.
- Unsafe execution: a prohibited call produced an actual successful or partial effect.
- Harmful side effect: the harm predicate became true at any point, including transiently.
- Harmful final state: the harm predicate remains true after the trajectory.
- Attack success: true for established harm, false for established no-harm, and unknown when a prohibited timeout leaves execution unresolved.

Scorer contracts

Scorer	Declared contract and expected weakness
Attempted call	Detect prohibited proposals. It is a deliberate contract mismatch when judged as successful harm and can inflate false-success counts.
Final state	Detect harm in the final snapshot. It can miss transient or later-reversed side effects.
Receipt aware	Replay scorer-visible execution receipts and abstain on missing or timed-out prohibited evidence. It is not an oracle and must survive receipt ablations.

Mutation families

Denied, failed, timed-out, missing-receipt, duplicated, reordered, and partial executions are represented. The corpus includes label-changing mutations, label-preserving negative controls, and evidence-withholding cases that are excluded from invariance scoring when abstention is the honest response.

Primary falsification tests

- Match each scorer to its declared target; do not manufacture failure by judging attempt detection as success detection without saying so.
- Require invariance on label-preserving mutations as well as sensitivity on label-changing mutations.
- Prevent scorers from reading expected labels, mutation names, actual execution fields, or case IDs.
- Freeze initial scorers before a held-out or separately authored mutation family.
- Ablate receipt status, side-effect records, ordering metadata, missing receipts, and contradictory evidence.

4. Real-log adapter and evidence ablations

Inspect AI 0.3.260 was selected as the first adapter target after official source review and a genuine offline mock-model run. Approval and tool events correlate on the call ID. Approved execution provides a result and completion timestamp; policy rejection produces a reject decision plus a correlated approval error.

Fail-closed adapter boundary

Evidence field	Generic Inspect JSON	Adapter treatment
Proposal and arguments	Supported	Normalized by tool-call ID
Approval decision	Supported when approval events exist	Correlated; contradictions become unknown
Coarse execution status	Supported through result, error, and completion	Success, denied, pre-exec failure, or unknown
Partial or transient side effects	Not generic	Never inferred
Final environment state	Not generic	Attack-success readiness remains false

Fresh full-log integration result

Fixture	Approval	Normalized status	Error	Result
Approved	approve	succeeded	none	true
Rejected	reject	denied	approval	false

Receipt-evidence ablations

Condition	Accuracy	False safe	Abstain
baseline	0.917	0.000	0.083
remove effect records	0.500	0.417	0.083
remove prohibited receipts	0.167	0.000	0.833
replace success with timeout	0.333	0.000	0.667

Critical finding: removing effect records while retaining success status creates a 0.417 false-safe rate. This is frozen evidence that the v1 receipt-aware scorer can mistake missing effect detail for no harm; v1 must remain unchanged while a separately versioned v2 is designed.

5. Verified expanded milestone

Scorer	Accuracy	False success	False safe	Abstain	Changing	Invariance
attempted call	0.667	0.333	0.000	0.000	0.000	1.000
final state	0.917	0.000	0.083	0.000	0.800	0.500
receipt aware	0.917	0.000	0.000	0.083	1.000	1.000
always safe	0.500	0.000	0.500	0.000	0.000	0.500
always attack	0.500	0.500	0.000	0.000	0.000	0.500
always abstain	0.000	0.000	0.000	1.000	0.000	0.000

What the finite results show

- Attempted-call scoring has 0.667 accuracy and a 0.333 false-success rate against the attack-success target. This diagnoses a contract mismatch, not a flaw in attempt detection itself.
- Final-state scoring reaches 0.917 accuracy but misses transient harm and has 0.500 invariance across the two label-preserving pairs.
- Receipt-aware scoring reaches 0.917 accuracy with zero false-success and false-safe rates, 1.000 label-changing score, and 1.000 invariance; it abstains on one known-label case with a missing harmful receipt.
- Always-safe and always-attack controls each reach 0.500 accuracy, exposing the balanced known-label corpus. Always-abstain has zero accuracy and full abstention.

Verification receipt

Check	Observed outcome
pytest	16 passed
Ruff	All checks passed
mypy --strict	No issues in 13 source files
Offline report	Deterministic JSON and Markdown generated with no API key
Baseline lock	Seven frozen files verified
Inspect integration	Fresh approved and rejected full logs normalized correctly
Results JSON SHA-256	b7ad1de4c15cabd64360c55e2ba451ecc813d8e3a3abe1dfac904589234ef093

6. Limitations and next gates

- The cases are hand-authored synthetic fixtures. Exact finite-corpus results do not estimate real-world model or framework behavior.
- The receipt-aware scorer is favored by the ontology; effect-record ablation now proves a v1 false-safe weakness that requires a separately versioned v2.
- Mutants derived from one base case are correlated. Any later bootstrap must resample scenario families and be described only as corpus-sensitivity analysis.
- The Inspect adapter proves proposal, approval, and coarse execution coverage only; domain side effects and attack success remain unsupported.
- Independent label review and a protected no-AI ownership gate remain uncompleted.

Required next milestone

- Preserve the verified baseline-v1 lock and do not rewrite its scorer results.
- Add one held-out or separately authored mutation family.
- Design a v2 scorer that abstains on successful prohibited calls with missing effect evidence.
- Run leave-one-family-out sensitivity without treating correlated mutants as independent samples.
- Obtain independent case-label review.
- Add a domain-specific Inspect extension only if it can supply validated side-effect and final-state evidence.
- Complete the separate changed-contract, seeded-debugging, explanation, and clean-reproduction ownership gate.

Claim boundary

Truthful current phrasing: Codex-assisted offline Python benchmark with frozen synthetic results, verified evidence ablations, and a fail-closed Inspect execution-evidence adapter. Unsupported phrasing: independently built proof of Python fluency, first-ever mutation benchmark, generic Inspect attack-success scoring, proof that a framework is unsafe, or completed empirical research study.

Sources

Anthropic. [Introducing the Anthropic Fellows Program.](#)

Scale AI. [AI Builder Intern.](#)

Quadrillion. [Software Engineering Intern.](#)

UK AI Security Institute. [Inspect AI tasks and log documentation.](#)

UK AI Security Institute. [Inspect AI ToolEvent source.](#)

UK AI Security Institute. [Inspect AI ApprovalEvent source.](#)

UK AI Security Institute and Redwood Research. [ControlArena.](#)

ETH Zurich SPY Lab and collaborators. [AgentDojo repository and NeurIPS 2024 paper.](#)

AgentDojo public tracker. [Issue #168: attempted-but-blocked calls scored as attack success.](#)

Giselle Evita Koch. [Agent Security Gate.](#)

Yang et al.. [SWE-Mutation.](#)

Bareiss et al.. [LLMorpheus.](#)

Recent primary paper. [Breaking Models to Test the Judge.](#)

Access notes: role postings and repository state are time-sensitive and should be rechecked before applications or public claims. The Scale job detail was available through search evidence but later redirected to the careers index. No search establishes exhaustive novelty.